

Computational security

The big picture

- ◆ we formally defined and constructed a perfectly secure cipher
- ◆ this encryption has some drawbacks
 - ◆ e.g., it employs a very large key
- ◆ by Shannon's Theorem, such limitations are unavoidable



Now, what?

Our approach: Relax “perfectness”

Initial model / abstraction

- ◆ **perfect secrecy** / security requires that
 - ◆ absolutely **no information is leaked** about the plaintext
 - ◆ to adversaries that **unlimited computational power**

Refined model / abstraction

- ◆ consider a relaxed notion of security, called **computational security**, where an encryption scheme is (for all **practical** purposes) secure even if
 - ◆ **a tiny amount of information is leaked** about the plaintext (e.g., w/ prob. 2^{-60})
 - ◆ to adversaries with **bounded computational power** (e.g., attacker invests 200ys)

Computational security

- ◆ to be contrasted against information-theoretic security
- ◆ de facto way in which security is modeled in most cryptographic settings
- ◆ an integral part of modern cryptography w/ rigorous mathematical proofs
- ◆ entails two relaxations
 - ◆ **security is guaranteed against efficient adversaries**
 - ◆ if an attacker invests in sufficiently large resources, it may break the scheme's security
 - ◆ goal: make required resources larger than those available to any realistic attacker!
 - ◆ **adversaries can potentially succeed** (i.e., security guarantees are probabilistic)
 - ◆ with some small probability the scheme is breakable
 - ◆ goal: make success probability sufficiently small so that it can be practically ignored!

Towards a rigorous definition of computational security

- ◆ **Concrete** approach
 - ◆ bounds the maximum success probability of any (randomized) adversary running for some specified amount of time or investing a specified amount of resources
 - ◆ general security result: *“A scheme is (t,ε) -secure if any adversary \mathcal{A} , running for time at most t , succeeds in breaking the scheme with probability at most ε*
 - ◆ need to
 - ◆ define what it means for an adversary to “break” a scheme
 - ◆ specify precisely the resources (e.g., time in seconds using a particular computer, or CPU cycles in a particular available supercomputer architecture)

Examples

- ◆ almost optimal security guarantees
 - ◆ key length n , key space size $|K| = 2^n$
 - ◆ A running for time t (e.g., CPU cycles) succeeds w/ prob. at most $ct/2^n$
 - ◆ this corresponds to a brute-force type of attack (w/out preprocessing)
- ◆ parameter c models advanced computing methods
 - ◆ c is typically larger than 1: e.g., parallelism can be used, etc.
- ◆ if $c = 1$, $n = 60$, security is enough for attackers running a desktop computer
 - ◆ 4 GHz (4×10^9 cycles/sec), 2^{60} CPU cycles require about 9 years
 - ◆ however, “fastest” available computer runs w/ 2×10^{16} cycles/sec, i.e., in ~1min!
 - ◆ choosing $n=80$ is better: the supercomputer would still need ~2 years

Today's recommendations

- ◆ recommended security parameter is $n = 128$
- ◆ large difference between 2^{80} and 2^{128}
 - ◆ #seconds since Big Bang is $\sim 2^{58}$
- ◆ if probability of success (within 1 year of computation) is $1/2^{60}$
 - ◆ it is more likely that Alice and Bob are hit by lighting (and thus don't care much about the breached confidentiality)
 - ◆ an event happening once in 100 years corresponds to probability 2^{-30} of happening at a particular second
- ◆ limitations of the concrete approach
 - ◆ harder to achieve, careful interpretation is needed, what if \mathcal{A} runs for $2t$ or $t/2$?

An alternative (less quantitative) approach

- ◆ **Asymptotic** approach
 - ◆ again, a security parameter n is used (e.g., key length)
 - ◆ **efficient (or realistic or feasible) adversaries** are equated with probabilistic poly-time (PPT) algorithms that run for time that is a polynomial of n
 - ◆ **small probability of success** is equated with success probabilities that are asymptotically smaller than any inverse polynomial in n
 - ◆ general security result: *“A scheme is secure if any PPT adversary A succeeds in breaking the scheme with at most negligible probability”*

Negligible functions (to capture tiny likelihood)

Typically, they measure the probability of success (of an attacker)

- ◆ Intuitively: very small probability
 - ◆ negligible, can be ignored, it's more likely to be hit by asteroid...
 - ◆ approaches 0 faster than the inverse of any polynomial
 - ◆ notation: **negl**
- ◆ Formally
 - ◆ A function $\mu : \mathbb{N} \rightarrow \mathbb{R}^+$ is negligible if for every positive integer **c** there exists an integer **N** such that for all **n > N**, it holds that

$$\mu(n) < 1 / n^c$$

3 equivalent definitions of perfect EAV-security

1) a posteriori = a priori

For every $\mathcal{D}_{\mathcal{M}}$, $m \in \mathcal{M}$ and $c \in C$, for which $\Pr [C = c] > 0$, it holds that

$$\Pr[M = m \mid C = c] = \Pr[M = m]$$

2) C is independent of M

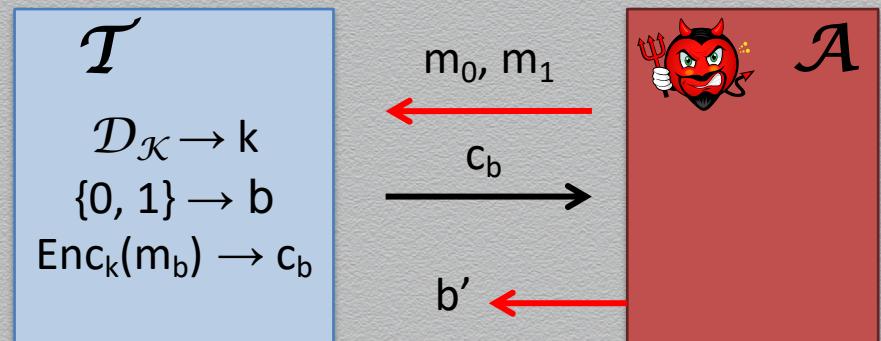
For every $m, m' \in \mathcal{M}$ and $c \in C$, it holds that

$$\Pr[\text{Enc}_k(m) = c] = \Pr[\text{Enc}_k(m') = c]$$

3) indistinguishability

For every \mathcal{A} , it holds that

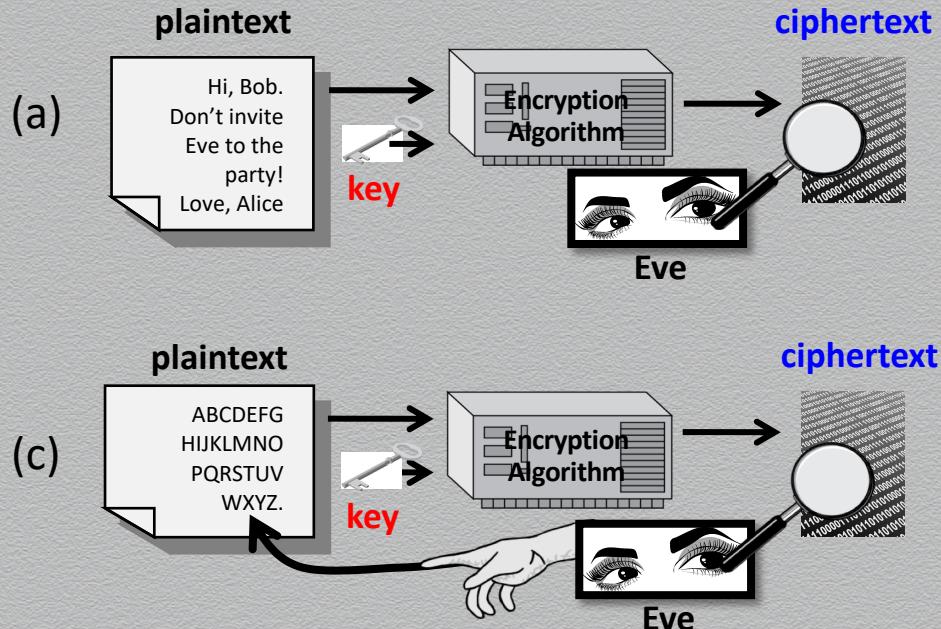
$$\Pr[b' = b] = 1/2$$



Recall: Possible eavesdropping attacks

An attacker may possess a

- ◆ (a) collection of ciphertexts
 - ◆ ciphertext only attack
 - ◆ **EAV-attack**
- ◆ (c) collection of plaintext/ciphertext pairs for plaintexts selected by the attacker
 - ◆ chosen plaintext attack
 - ◆ CPA-attack



From perfect to computational EAV-security

- ◆ **perfect** security: $M, \text{Enc}_K(M)$ are independent
 - ◆ absolutely **no information is leaked** about the plaintext
 - ◆ to adversaries that **unlimited computational power**
- ◆ **computational** security: for all **practical** purposes, $M, \text{Enc}_K(M)$ are independent
 - ◆ **a tiny amount of information is leaked** about the plaintext (e.g., w/ prob. 2^{-60})
 - ◆ to adversaries with **bounded computational power** (e.g., attacker invests 200ys)
- ◆ attacker's **best strategy** remains **ineffective**
 - ◆ **random guess** on secret key; or
 - ◆ **exhaustive search** over key space (**brute force attack**)

Relaxing indistinguishability

Relax the definition of perfect secrecy – that is based on indistinguishability

- ◆ require that m_0, m_1 are chosen by a **PPT adversary**
- ◆ require that no **PPT adversary** can distinguish $\text{Enc}_k(m_0)$ from $\text{Enc}_k(m_1)$

non-negligibly better than guessing

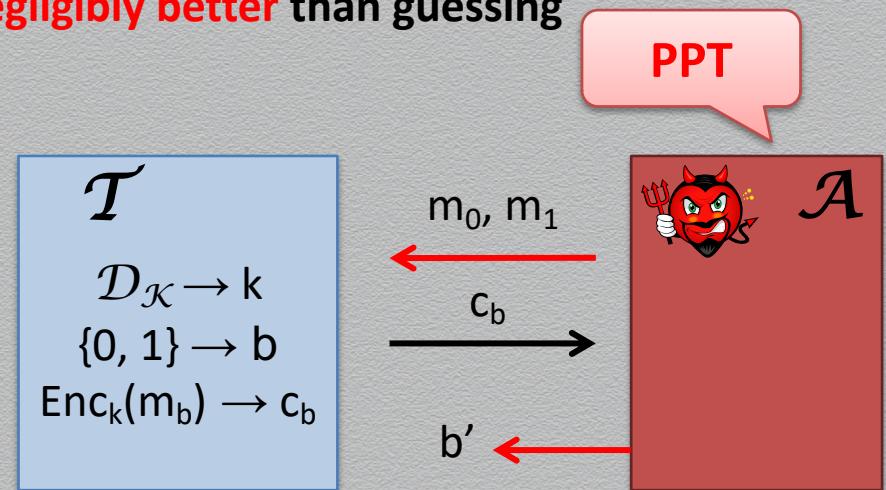
3) indistinguishability

For every \mathcal{A} , it holds that

$$\Pr[b' = b] = 1/2 + \text{negl}$$

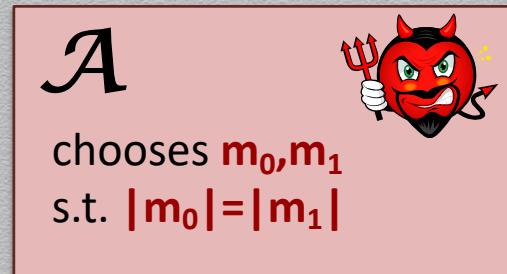
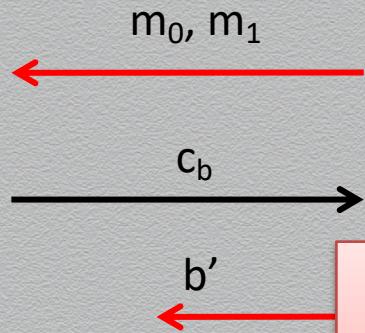
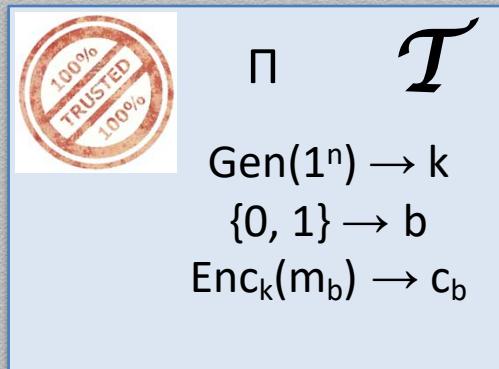
PPT

PPT



Game-based definition of computational EAV-security

encryption scheme $\Pi = \{\mathcal{M}, (\text{Gen}, \text{Enc}, \text{Dec})\}$



Alternatively:
"is semantically secure"

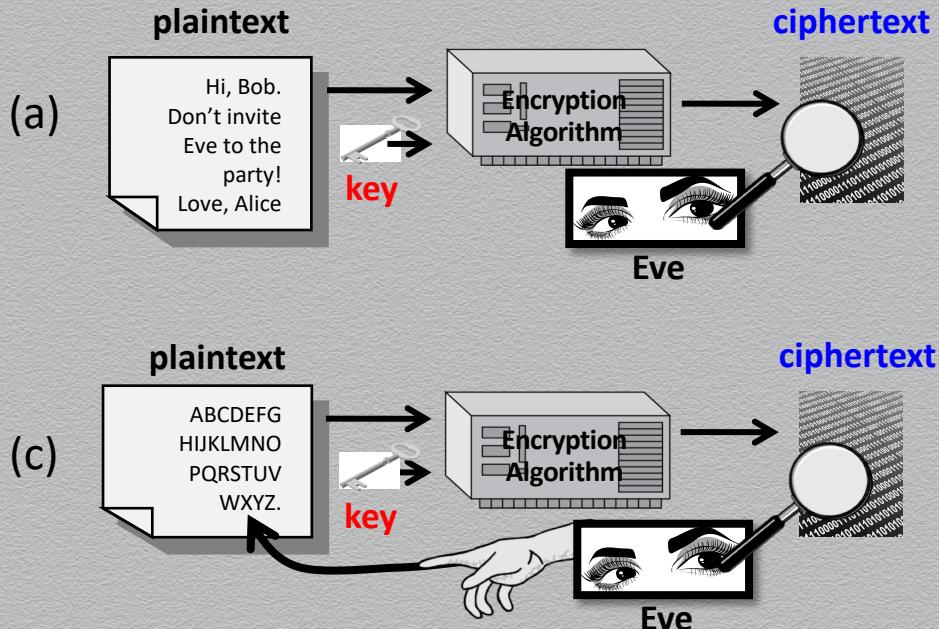
We say that (Enc, Dec) is **EAV-secure** if any PPT \mathcal{A} has probability at most $0.5 + \varepsilon(n)$, where ε is a negligible function

I.e., no PPT \mathcal{A} computes b correctly non-negligibly better than randomly guessing

Recall: Possible eavesdropping attacks

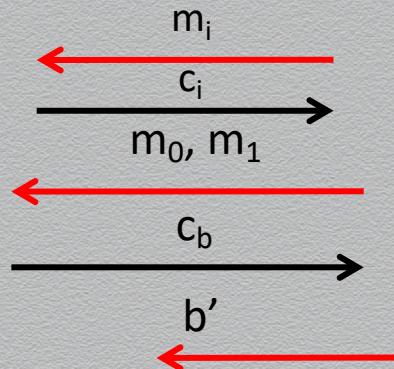
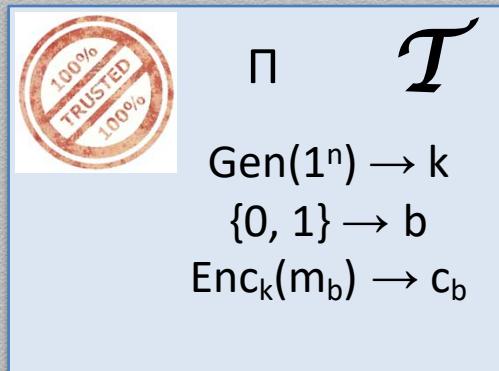
An attacker may possess a

- ◆ (a) collection of ciphertexts
 - ◆ ciphertext only attack
 - ◆ EAV-attack
- ◆ (c) collection of plaintext/ciphertext pairs for plaintexts selected by the attacker
 - ◆ chosen plaintext attack
 - ◆ CPA-attack



Similarly: CPA-security

encryption scheme $\Pi = \{\mathcal{M}, (\text{Gen}, \text{Enc}, \text{Dec})\}$



We say that (Enc, Dec) is **CPA-secure** if any PPT adversary guesses b correctly with probability at most $0.5 + \varepsilon(n)$, where ε is a negligible function

I.e., no PPT \mathcal{A} computes b correctly non-negligibly better than randomly guessing,
even when it learns the encryptions of messages of its choice

On CPA security

Facts

- ◆ Any encryption scheme that is CPA-secure is also CPA-secure for multiple encryptions
- ◆ **CPA security implies probabilistic encryption – can you see why?**
- ◆ EAV-security for multiple messages implies probabilistic encryption